

CSC 4020 Fundamentals of Machine Learning: Principal Component Analysis I

Baoyuan Wu

May 8/10

[Slide credit: Mengye Ren, Matthew MacKay]

Overview

Today we'll cover the first unsupervised learning algorithm for this course: principal component analysis (PCA)

Dimensionality reduction: map the data to a lower dimensional space

- Save computation/memory

- Reduce overfitting

- Visualize in 2 dimensions

PCA is a linear model, with a closed-form solution. It's useful for understanding lots of other algorithms.

- Autoencoders

- Matrix factorizations (next lecture)

Projection onto a subspace

Set-up: given a dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\} \subset \mathbb{R}^D$

Set $\boldsymbol{\mu}$ to the mean of the data, $\boldsymbol{\mu} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}^{(i)}$

Goal: find a K -dimensional subspace $\mathcal{S} \subset \mathbb{R}^D$ such that $\mathbf{x}^{(n)} - \boldsymbol{\mu}$ is “well-represented” by its projection onto \mathcal{S}

Recall: The **projection** of a point \mathbf{x} onto \mathcal{S} is the point in \mathcal{S} closest to \mathbf{x} .

Projection onto a subspace

Let $\{\mathbf{u}_k\}_{k=1}^K$ be an orthonormal basis of the subspace \mathcal{S}

Approximate each data point \mathbf{x} as:


$$\begin{aligned}\tilde{\mathbf{x}} &= \boldsymbol{\mu} + \text{Proj}_{\mathcal{S}}(\mathbf{x} - \boldsymbol{\mu}) \\ &= \boldsymbol{\mu} + \sum_{k=1}^K z_k \mathbf{u}_k\end{aligned}$$

From linear algebra: $z_k = \mathbf{u}_k^T(\mathbf{x} - \boldsymbol{\mu})$

Let \mathbf{U} be a matrix with columns $\{\mathbf{u}_k\}_{k=1}^K$ then $\mathbf{z} = \mathbf{U}^T(\mathbf{x} - \boldsymbol{\mu})$

Also: $\tilde{\mathbf{x}} = \boldsymbol{\mu} + \mathbf{U}\mathbf{z}$

Projection onto a Subspace



$$\mathbf{z} = \mathbf{U}^\top (\mathbf{x} - \mu)$$

In machine learning, $\tilde{\mathbf{x}}$ is also called the **reconstruction** of \mathbf{x} .
 \mathbf{z} is its **representation**, or **code**.

Projection onto a Subspace


If we have a K -dimensional subspace in a D -dimensional input space, then $\mathbf{x} \in \mathbb{R}^D$ and $\mathbf{z} \in \mathbb{R}^K$.

If the data points \mathbf{x} all lie close to their reconstructions, then we can approximate distances, etc. in terms of these same operations on the code vectors \mathbf{z} .

If $K \ll D$, then it's much cheaper to work with \mathbf{z} than \mathbf{x} .

A mapping to a space that's easier to manipulate or visualize is called a [representation](#), and learning such a mapping is [representation learning](#).

Mapping data to a low-dimensional space is called [dimensionality reduction](#).



Learning a Subspace

How to choose a good subspace \mathcal{S} ?

Need to choose $D \times K$ matrix \mathbf{U} with orthonormal columns.

Two criteria:

Minimize the **reconstruction error**

$$\min \frac{1}{N} \sum_{i=1}^N \|\mathbf{x}^{(i)} - \tilde{\mathbf{x}}^{(i)}\|^2$$

Maximize the variance of the code vectors

$$\begin{aligned} \max_j \sum_j \text{Var}(z_j) &= \frac{1}{N} \sum_j \sum_i (z_j^{(i)} - \bar{z}_j)^2 \\ &= \frac{1}{N} \sum_i \|\mathbf{z}^{(i)} - \bar{\mathbf{z}}\|^2 \\ &= \frac{1}{N} \sum_i \|\mathbf{z}^{(i)}\|^2 \end{aligned}$$

Exercise: show $\bar{\mathbf{z}} = \mathbf{0}$

Note: here, $\bar{\mathbf{z}}$ denotes the mean, not a derivative.

Learning a Subspace


These two criteria are equivalent! I.e., we'll show

$$\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}^{(i)} - \tilde{\mathbf{x}}^{(i)}\|^2 = \text{const} - \frac{1}{N} \sum_i \|\mathbf{z}^{(i)}\|^2$$

Observation: by unitarity,

$$\|\tilde{\mathbf{x}}^{(i)} - \boldsymbol{\mu}\| = \|\mathbf{U}\mathbf{z}^{(i)}\| = \|\mathbf{z}^{(i)}\|$$

By the Pythagorean Theorem,



$$\underbrace{\frac{1}{N} \sum_{i=1}^N \|\tilde{\mathbf{x}}^{(i)} - \boldsymbol{\mu}\|^2}_{\text{projected variance}} + \underbrace{\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}^{(i)} - \tilde{\mathbf{x}}^{(i)}\|^2}_{\text{reconstruction error}} \\ = \underbrace{\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}^{(i)} - \boldsymbol{\mu}\|^2}_{\text{constant}}$$

Principal Component Analysis

Choosing a subspace to maximize the projected variance, or minimize the reconstruction error, is called **principal component analysis (PCA)**.

Recall:

Spectral Decomposition: a symmetric matrix \mathbf{A} has a full set of eigenvectors, which can be chosen to be orthogonal. This gives a decomposition

$$\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T,$$

where \mathbf{Q} is orthogonal and $\mathbf{\Lambda}$ is diagonal. The columns of \mathbf{Q} are eigenvectors, and the diagonal entries λ_j of $\mathbf{\Lambda}$ are the corresponding eigenvalues.

I.e., symmetric matrices are diagonal in some basis.

A symmetric matrix \mathbf{A} is positive semidefinite iff each $\lambda_j \geq 0$.

Principal Component Analysis

Consider the **empirical covariance matrix**:

$$\mathbf{\Sigma} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \boldsymbol{\mu})(\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top$$


Recall: Covariance matrices are symmetric and positive semidefinite.

The optimal PCA subspace is spanned by the top K eigenvectors of $\mathbf{\Sigma}$.

More precisely, choose the first K of any orthonormal eigenbasis for $\mathbf{\Sigma}$.

The general case is tricky, but we'll show this for $K = 1$.

These eigenvectors are called **principal components**, analogous to the principal axes of an ellipse.



Deriving PCA

For $K = 1$, we are fitting a unit vector \mathbf{u} , and the code is a scalar $z = \mathbf{u}^\top (\mathbf{x} - \boldsymbol{\mu})$.

$$\begin{aligned}\frac{1}{N} \sum_i [z^{(i)}]^2 &= \frac{1}{N} \sum_i (\mathbf{u}^\top (\mathbf{x}^{(i)} - \boldsymbol{\mu}))^2 \\&= \frac{1}{N} \sum_{i=1}^N \mathbf{u}^\top (\mathbf{x}^{(i)} - \boldsymbol{\mu}) (\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top \mathbf{u} \\&= \mathbf{u}^\top \left[\frac{1}{N} \sum_{i=1}^N (\mathbf{x}^{(i)} - \boldsymbol{\mu}) (\mathbf{x}^{(i)} - \boldsymbol{\mu})^\top \right] \mathbf{u} \\&= \mathbf{u}^\top \boldsymbol{\Sigma} \mathbf{u} \\&= \mathbf{u}^\top \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^\top \mathbf{u} \\&= \mathbf{a}^\top \boldsymbol{\Lambda} \mathbf{a} \\&= \sum_{j=1}^D \lambda_j a_j^2\end{aligned}$$

Spectral Decomposition
for $\mathbf{a} = \mathbf{Q}^\top \mathbf{u}$

Deriving PCA

Maximize $\mathbf{a}^\top \mathbf{\Lambda} \mathbf{a} = \sum_{j=1}^D \lambda_j a_j^2$ for $\mathbf{a} = \mathbf{Q}^\top \mathbf{u}$.

This is a change-of-basis to the eigenbasis of $\mathbf{\Sigma}$.

Assume the λ_i are in sorted order. For simplicity, assume they are all distinct.

Observation: since \mathbf{u} is a unit vector, then by unitarity, \mathbf{a} is also a unit vector. I.e., $\sum_j a_j^2 = 1$.

By inspection, set $a_1 = \pm 1$ and $a_j = 0$ for $j \neq 1$.

Hence, $\mathbf{u} = \mathbf{Q} \mathbf{a} = \mathbf{q}_1$ (the top eigenvector).

A similar argument shows that the k th principal component is the k th eigenvector of $\mathbf{\Sigma}$. If you're interested, look up the [Courant-Fischer Theorem](#).

Decorrelation

Interesting fact: the dimensions of \mathbf{z} are decorrelated. For now, let Cov denote the empirical covariance.

$$\begin{aligned}\text{Cov}(\mathbf{z}) &= \text{Cov}(\mathbf{U}^\top (\mathbf{x} - \boldsymbol{\mu})) \\ &= \mathbf{U}^\top \text{Cov}(\mathbf{x}) \mathbf{U} \\ &= \mathbf{U}^\top \boldsymbol{\Sigma} \mathbf{U} \\ &= \mathbf{U}^\top \mathbf{Q} \boldsymbol{\Lambda} \mathbf{Q}^\top \mathbf{U} \\ &= \begin{pmatrix} \mathbf{I} & \mathbf{0} \end{pmatrix} \boldsymbol{\Lambda} \begin{pmatrix} \mathbf{I} \\ \mathbf{0} \end{pmatrix} && \text{by orthogonality} \\ &= \text{top left } K \times K \text{ block of } \boldsymbol{\Lambda}\end{aligned}$$

If the covariance matrix is diagonal, this means the features are uncorrelated.

This is why PCA was originally invented (in 1901!).

Recap

Recap:

Dimensionality reduction aims to find a low-dimensional representation of the data.

PCA projects the data onto a subspace which maximizes the projected variance, or equivalently, minimizes the reconstruction error.

The optimal subspace is given by the top eigenvectors of the empirical covariance matrix.

PCA gives a set of decorrelated features.

Applying PCA to faces

Consider running PCA on 2429 19x19 grayscale images (CBCL data)

Can get good reconstructions with only 3 components



PCA for pre-processing: can apply classifier to latent representation

For face recognition PCA with 3 components obtains 79% accuracy on face/non-face discrimination on test data vs. 76.8% for a Gaussian mixture model (GMM) with 84 states. (We'll cover GMMs later in the course.)

Can also be good for visualization

Applying PCA to faces: Learned basis

Principal components of face images (“eigenfaces”)



Applying PCA to digits



reconstructed with 2 bases



reconstructed with 10 bases



reconstructed with 100 bases



reconstructed with 506 bases



mean



principal basis 1



principal basis 2



principal basis 3



Next: two more interpretations of PCA, which have interesting generalizations.


- Autoencoders

- Matrix factorization (later lecture)

Autoencoders

An **autoencoder** is a feed-forward neural net whose job it is to take an input \mathbf{x} and predict \mathbf{x} .

To make this non-trivial, we need to add a **bottleneck layer** whose dimension is much smaller than the input.



Linear Autoencoders

Why autoencoders?

- Map high-dimensional data to two dimensions for visualization

- Learn abstract features in an unsupervised way so you can apply them to a supervised task

 - Unlabeled data can be much more plentiful than labeled data

Linear Autoencoders

The simplest kind of autoencoder has one hidden layer, linear activations, and squared error loss.


$$\mathcal{L}(\mathbf{x}, \tilde{\mathbf{x}}) = \|\mathbf{x} - \tilde{\mathbf{x}}\|^2$$

This network computes $\tilde{\mathbf{x}} = \mathbf{W}_2 \mathbf{W}_1 \mathbf{x}$, which is a linear function.

If $K \geq D$, we can choose \mathbf{W}_2 and \mathbf{W}_1 such that $\mathbf{W}_2 \mathbf{W}_1$ is the identity matrix. This isn't very interesting.

But suppose $K < D$:

\mathbf{W}_1 maps \mathbf{x} to a K -dimensional space, so it's doing dimensionality reduction.




Linear Autoencoders

Observe that the output of the autoencoder must lie in a K -dimensional subspace spanned by the columns of \mathbf{W}_2 .

We saw that the best possible K -dimensional subspace in terms of reconstruction error is the PCA subspace.

The autoencoder can achieve this by setting $\mathbf{W}_1 = \mathbf{U}^\top$ and $\mathbf{W}_2 = \mathbf{U}$.

Therefore, the optimal weights for a linear autoencoder are just the principal components!





Nonlinear Autoencoders

Deep nonlinear autoencoders learn to project the data, not onto a subspace, but onto a nonlinear **manifold**


This manifold is the image of the decoder.

This is a kind of **nonlinear dimensionality reduction**.



Nonlinear Autoencoders

Nonlinear autoencoders can learn more powerful codes for a given dimensionality, compared with linear autoencoders (PCA)



Nonlinear Autoencoders

Here's a 2-dimensional autoencoder representation of newsgroup articles. They're color-coded by topic, but the algorithm wasn't given the labels.

