



MAT 3007 – Optimization

Algorithms for Unconstrained Optimization Problems

Lecture 16

July 14th

Andre Milzarek

SDS / CUHK-SZ



Repetition



Convex Problems:

- ▶ A minimization problem $\min_{x \in \Omega} f(x)$ is called **convex** if Ω is a convex set and f is convex.
- ▶ Convexity/concavity plays a very important role in optimization problems!

Calculus & Rules:

- ▶ A function f is convex on a convex set Ω iff the Hessian $\nabla^2 f(x)$ is positive semidefinite for all $x \in \Omega$.
- ▶ Rich calculus: sum rule, composition, max-/min-rule.
- ▶ If f is convex, then $L_{\leq c} = \{x : f(x) \leq c\}$ is a convex set \rightsquigarrow can be used to check convexity of constraints.
- ▶ $\Omega = \{x : g(x) = 0, h(x) = 0\}$ is convex if all g_i are convex and h is an **affine-linear** function, i.e., $h(x) = Ax - b$.



Convexity & Optimality

- ▶ Every local minimizer of a convex problem is a **global minimizer**.
- ▶ Every stationary point or KKT point of a (unconstrained/constrained) convex problem is a **global minimizer**.
- ↪ If f is **concave**, we typically consider $\max_{x \in \Omega} f(x)$ or $\min_{x \in \Omega} -f(x)$.



Algorithms for Unconstrained Problems

We start with the unconstrained problem:

$$\text{minimize}_{x \in \mathbb{R}^n} f(x)$$

We are going to study the following methods:

- ▶ Bisection search and golden section search.
- ▶ Gradient descent method.
- ▶ Newton's method.

Optimization algorithms are **iterative procedures**:

- ▶ Starting from an **initial point** x^0 , a sequence of iterates $\{x^k\}$ is generated.
- ▶ **Goal**: reduction of the function values and convergence to an optimal solution.



Problems in \mathbb{R}



Assume $f : \mathbb{R} \rightarrow \mathbb{R}$ is a single variable function.

Our Objective: find a local minimizer of f .


We introduce two methods:

- ▶ Bisection method.
- ▶ Golden section method.

Bisection method uses the idea that the local minimizer must satisfy the first-order necessary conditions: $f'(x) = 0$.

Therefore, the problem becomes a root-finding problem for

$$g(x) = f'(x) = 0.$$



Assume we can find x_ℓ and x_r such that $g(x_\ell) < 0$ and $g(x_r) > 0$.

By the **intermediate value theorem**, if g is continuous, there must exist a root of g in $[x_\ell, x_r]$.

Bisection Method

1. Define $x_m = \frac{x_\ell + x_r}{2}$.
2. If $g(x_m) = 0$, then output x_m .
3. Otherwise:
 - If $g(x_m) > 0$, then let $x_r = x_m$.
 - If $g(x_m) < 0$, then let $x_\ell = x_m$.
4. If $|x_r - x_\ell| < \epsilon$: stop and output $\frac{x_\ell + x_r}{2}$, otherwise go back to step 1.

One can also set the stopping criterion based on $|g(x)| < \epsilon$.



In the bisection method, each iteration will divide the search interval to half.

Therefore, to find an ϵ approximation of x^* , we need at most $\log_2 \frac{x_r - x_\ell}{\epsilon}$ many iterations.

Applying the bisection method to f' , we can find an approximate stationary point. If f is convex, this is an (approximate) global minimizer of f .

- ▶ Although simple, the bisection method is very useful in practice because it is easy to implement.

Example: Use bisection method to minimize:

$$f(x) = -\frac{xe^{-x}}{1 + e^{-x}} \quad \rightsquigarrow \quad f'(x) = -\frac{e^{-x}(1 - x + e^{-x})}{(1 + e^{-x})^2}$$

```
1 function [x,gx] = bisection(g,xl,xr,options)
2
3 % Compute initial function values
4 gr = g(xr); gl = g(xl); sl = sign(gl);
5
6 if gl*gr > 0
7     fprintf(1,'The input data not suitable!');
8     x = []; gx = []; return
9 end
10
11 for i = 1:options.maxit
12     xm = (xl + xr)/2; gm = g(xm);
13
14     if abs(gm) < options.tol || abs(xl-xr) < options.tol
15         x = xm; gx = gm; return
16     end
17
18     if gm > 0
19         if sl < 0, xr = xm; else, xl = xm; end
20     else
21         if sl < 0, xl = xm; else, xr = xm; end
22     end
23 end
```



Drawback of the bisection method: When solving (single variable, unconstrained) optimization problems, we require the knowledge (and computation) of f' .

- Sometimes, f' is not available. For example, f sometimes is only a **black box**, which does not admit an analytical form (thus, the derivative is hard to compute)


However, if we know that f has a unique local minimum x^* in the range $[x_\ell, x_r]$, then we still have a very efficient way to find x^* :

- We call f **unimodal** if it only has one single stationary point (on \mathbb{R}).
- Unimodal functions have the property that the local minimum is already global. (Similarly, if the stationary point is a local maximum).

Example of a Unimodal Function



Consider $f(x) = -\frac{xe^{-x}}{1+e^{-x}}$:



This is a unimodal function, but not a concave function.

Golden Section Method

Assume we start with $[x_\ell, x_r]$. Assume $0 < \phi < 0.5$.


1. Set $x'_\ell = \phi x_r + (1 - \phi)x_\ell$ and $x'_r = (1 - \phi)x_r + \phi x_\ell$.
2. If $f(x'_\ell) < f(x'_r)$, then the minimizer must lie in $[x_\ell, x'_r]$, so set $x_r = x'_r$.
3. Otherwise, the minimizer must lie in $[x'_\ell, x_r]$, so set $x_\ell = x'_\ell$.
4. If $x_r - x_\ell < \epsilon$, output $\frac{x_\ell + x_r}{2}$, otherwise go back to step 1.

► Suppose we update $x_r = x'_r$. We want to choose ϕ such that x'_r of the new iteration coincides with x'_ℓ of the old iteration.

~> This allows to save one function evaluation!

► This is true when

$$\phi = \frac{3 - \sqrt{5}}{2} \quad \text{and} \quad 1 - \phi = \frac{\sqrt{5} - 1}{2} = 0.618.$$



Both the bisection and golden section method can be easily adapted for maximization problems. (Just adjust the comparison).

Example Revisited: Use the Golden section method to maximize:

$$f(x) = \frac{xe^{-x}}{1 + e^{-x}}$$

Higher-Dimensional Problems

Next, we consider the n -dimensional problem:

$$\text{minimize}_{x \in \mathbb{R}^n} f(x)$$

- ▶ There is no clear bisection or golden section in that case.

Solution and General Idea:

- ▶ Each time, we first find a **search direction**.
- ▶ Then, we search for a good next step along that direction (which reduces to a one-dimensional problem).



Starting from the **initial point** x^0 , we generate a sequence of points:

$$x^{k+1} = x^k + \alpha_k d^k.$$

We call d^k the **search direction** (a vector) and α_k the **step size** (a scalar).

- ▶ The key is to choose a proper direction d^k at each iteration.
- ▶ d^k typically depends on x^k .
- ▶ The step size α_k may be chosen in accordance with some line (one-dimensional) search rules (later).

We will study two such methods:

- ▶ Gradient descent method and Newton's method.

In the following, we assume that f is continuously differentiable.

Definition: Descent Direction

A vector $d \in \mathbb{R}^n$ is a **descent direction** of f at x if $\nabla f(x)^\top d < 0$.

Important Observation:

- ▶ Taking a small enough step along a descent direction reduces the objective function value.
- ▶ By Taylor: there exists $\epsilon > 0$ such that

$$f(x + \alpha d) < f(x) \quad \forall \alpha \in (0, \epsilon].$$



Schematic Descent Directions Method

1. Initialization: Select an initial point $x^0 \in \mathbb{R}^n$.

For $k = 0, 1, \dots$:

2. Pick a **descent direction** d^k .
3. Find a **stepsize** α_k satisfying $f(x^k + \alpha_k d^k) < f(x^k)$.
4. Set $x^{k+1} = x^k + \alpha_k d^k$.
5. If a **stopping criterion** is satisfied, then STOP and x^{k+1} is the output.

Open questions and missing details:

- ▶ What is the initial point x^0 ?
- ▶ How to choose the descent direction? What step size should be taken?
- ▶ What is the stopping criterion?



Gradient Descent:

- ▶ One simple and possible descent direction is $d^k = -\nabla f(x^k)$. This direction satisfies:

$$\nabla f(x^k)^\top d^k = -\|\nabla f(x^k)\|^2 < 0$$

as long as $\nabla f(x^k) \neq 0$.

- ▶ Choosing $d^k = -\nabla f(x^k)$, the abstract descent method becomes the **gradient descent method**.

Stopping Criterion:

- ▶ A popular stopping criterion is: $\|\nabla f(x^{k+1})\| \leq \epsilon$ with **tolerance** $\epsilon > 0$.
- ↪ We stop if x^{k+1} is an **approximate stationary point**.



Constant Step Size:

- ▶ Choose $\alpha_k = \bar{\alpha}$ for all k .

Exact Line Search:

- ▶ An intuitive idea is to choose α_k to achieve the largest descent

That is, choose α_k such that:

$$\alpha_k = \operatorname{argmin}_{\alpha \geq 0} f(x^k + \alpha d^k). \quad (1)$$

- ▶ If we get the exact α_k in (1), we say we used an **exact line search** method to find the step size.
- ▶ We can use the golden section method to perform the exact line search.
- ▶ In some situations, we can even find the exact α analytically.

Consider

$$f(x) = b^\top x + \frac{1}{2}x^\top Ax \quad (A \text{ positive definite})$$

At x^k , the gradient descent method will choose:

$$d^k = -\nabla f(x^k) = -(b + Ax^k).$$

To choose the step size, notice that we can explicitly compute

$$\begin{aligned} f(x^k + \alpha d^k) &= b^\top (x^k + \alpha d^k) + \frac{1}{2}(x^k + \alpha d^k)^\top A(x^k + \alpha d^k) \\ &= \frac{1}{2}\alpha^2 (d^k)^\top A d^k + \alpha(b^\top d^k + (x^k)^\top A d^k) + f(x^k) \end{aligned}$$

This is a quadratic function of α with positive second-order term!
We can find the optimal $\alpha \geq 0$ minimizing $\phi(\alpha) = f(x^k + \alpha d^k)$:

$$\alpha_k = \frac{(d^k)^\top d^k}{(d^k)^\top A d^k}.$$

Example: Gradient Method for Quadratic Functions



```
1 function [x,obj] = gm_quadratic(A,b,x0,eps)
2
3 x      = x0;      iter    = 0;
4 g      = A*x + b; ng      = norm(g);
5
6 fprintf(1,'--- grad. method ; n = %g\n',length(b));
7 fprintf(1,'ITER ; OBJ.VAL ; G.NORM ; STEP.SIZE\n');
8
9 while ng > eps && iter < 10000
10     iter    = iter + 1;
11     alpha   = ng^2 / (g'*A*g);
12     x       = x - alpha*g;
13     g       = A*x + b;
14     ng      = norm(g);
15     obj     = 0.5*x'*A*x + b'*x;
16     fprintf(1,'[%4i] ; %2.6f ; %2.6f ; %1.2f\n',iter,obj,ng,
17             alpha);
18 end
```



We now want to test the method and solve the problem:

$$\min_x f(x) = x_1^2 + 2x_2^2 = \frac{1}{2}x^\top \begin{pmatrix} 2 & 0 \\ 0 & 4 \end{pmatrix} x.$$

We use the initial point $x^0 = (2, 1)^\top$ and the tolerance $\varepsilon = 10^{-5}$.

The method stops after 13 iterations with a solution that is already very close the optimal value $x = 10^{-5} \cdot (0.1254, -0.0627)^\top$.



In general, we can not expect that

$$\alpha_k = \operatorname{argmin}_{\alpha \geq 0} f(x^k + \alpha d^k) \quad (2)$$

can be solved explicitly. It can be very time-consuming!

- ▶ Computing α_k is an optimization problem on its own!
- ▶ It is also not clear how much benefit there is when solving (2) exactly. After all, it is just one iteration and it does not imply that $x^k + \alpha_k d^k$ is optimal.

Agenda: Let us consider approximate and cheaper techniques!

- ▶ There are multiple ways to do it, here we introduce the **backtracking line search** technique.



Assume we have found a descent direction d^k and we want to choose step size α_k .

Let $\sigma, \gamma \in (0, 1)$ be given. Choose α_k as the largest element in $\{1, \sigma, \sigma^2, \sigma^3, \dots\}$ such that

$$f(x^k + \alpha_k d^k) - f(x^k) \leq \gamma \alpha_k \cdot \nabla f(x^k)^\top d^k.$$

- ▶ This condition is called **Armijo condition**.
- ▶ α_k can be determined after finitely many steps if d^k is a **descent direction**.

Procedure:

1. Start with $\alpha = 1$.
2. If $f(x^k + \alpha d^k) \leq f(x^k) + \gamma \alpha \cdot \nabla f(x^k)^\top d^k$, choose $\alpha_k = \alpha$. Otherwise, set $\alpha = \sigma \alpha$ and repeat this step.



Why does this work?

- By Taylor expansion, if α is sufficiently small, we have

$$f(x^k + \alpha d^k) \approx f(x^k) + \alpha \nabla f(x^k)^\top d^k < f(x^k) + \gamma \alpha \cdot \nabla f(x^k)^\top d^k.$$

Therefore, as long as α is small enough, the Armijo condition must be satisfied (recall $\nabla f(x^k)^\top d^k = -\|\nabla f(x^k)\|^2 < 0$).

Illustration:

- Define $\phi_k(\alpha) := f(x^k + \alpha d^k) - f(x^k)$. Then, we have

$$\phi'_k(\alpha) = \nabla f(x^k + \alpha d^k)^\top d^k, \quad \phi'_k(0) = \nabla f(x^k)^\top d^k.$$

- The Armijo condition is then equivalent to:

$$\text{find } \alpha \text{ with } \phi_k(\alpha) \leq \gamma \alpha \cdot \phi'_k(0).$$

- Notice that $\phi'_k(0) < 0$ (since d^k is a descent direction).



Gradient Descent Method

1. Initialization: Select an initial point $x^0 \in \mathbb{R}^n$.

For $k = 0, 1, \dots$:

2. Pick a **stepsize** α^k by a line search procedure (exact line search or backtracking) on the function


$$\phi(\alpha) = f(x^k - \alpha \nabla f(x^k)).$$

3. Set $x^{k+1} = x^k - \alpha_k \nabla f(x^k)$.
4. If $\|\nabla f(x^{k+1})\| \leq \varepsilon$, then STOP and x^{k+1} is the output.

Minimize

$$f(x) = \exp(x_1 + x_2) + x_1^2 + 3x_2^2 - x_1x_2$$

using the gradient method with Armijo line search.



Questions?